

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from matplotlib import pyplot as plt
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
df = pd.read_csv('/content/drive/MyDrive/archive
(8)/iphone_purchase_records.csv')
```

```
df.head()
```

```
{"summary":{"\n  \"name\": \"df\",\n  \"rows\": 400,\n  \"fields\": [\n    {\n      \"column\": \"Gender\",\n      \"properties\": {\n        \"dtype\": \"category\",\n        \"num_unique_values\": 2,\n        \"samples\": [\n          \"Female\",\n          \"Male\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Age\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 10,\n        \"min\": 18,\n        \"max\": 60,\n        \"num_unique_values\": 43,\n        \"samples\": [\n          50,\n          39\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Salary\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 34096,\n        \"min\": 15000,\n        \"max\": 150000,\n        \"num_unique_values\": 117,\n        \"samples\": [\n          117000,\n          76000\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Purchase Iphone\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0,\n        \"min\": 0,\n        \"max\": 1,\n        \"num_unique_values\": 2,\n        \"samples\": [\n          1,\n          0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ],\n  \"type\": \"dataframe\", \"variable_name\": \"df\"}
```

```
df.describe()
```

```
{"summary":{"\n  \"name\": \"df\",\n  \"rows\": 8,\n  \"fields\": [\n    {\n      \"column\": \"Age\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 130.27423677374767,\n        \"min\": 10.482876597307914,\n        \"max\": 400.0,\n        \"num_unique_values\": 8,\n        \"samples\": [\n          37.655,\n          37.0,\n          400.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Salary\",\n      \"properties\": {\n        \"dtype\": \"number\",
```

```

\"std\": 47214.004060407126,\n        \"min\": 400.0,\n\"max\": 150000.0,\n        \"num_unique_values\": 8,\n\"samples\": [\n        69742.5,\n        70000.0,\n        400.0\n    ],\n        \"semantic_type\": \"\",\n\"description\": \"\"\n    }\n    },\n    {\n        \"column\":\n        \"Purchase Iphone\",\n        \"properties\": {\n            \"dtype\":\n            \"number\",\n            \"std\": 141.27865845809384,\n            \"min\":\n            0.0,\n            \"max\": 400.0,\n            \"num_unique_values\": 5,\n            \"samples\": [\n                0.3575,\n                1.0,\n                0.479863963596869\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\"\n        }\n    }\n    ]\n}], \"type\": \"dataframe\"}

```

```
df.isnull().sum()
```

```

Gender      0
Age         0
Salary      0
Purchase Iphone  0
dtype: int64

```

```
df = pd.get_dummies(df, columns=['Gender'], drop_first=True)
```

```
X = df[['Age', 'Salary']] # Select 'Age' and 'Salary' as independent variables
```

```
y = df['Purchase Iphone']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.25, random_state=0)
```

```
!pip install scikit-learn
```

```
from sklearn.preprocessing import StandardScaler # Import StandardScaler
```

```
sc = StandardScaler()
```

```
X_train = sc.fit_transform(X_train)
```

```
X_test = sc.transform(X_test)
```

```
Requirement already satisfied: scikit-learn in
```

```
/usr/local/lib/python3.10/dist-packages (1.2.2)
```

```
Requirement already satisfied: numpy>=1.17.3 in
```

```
/usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.25.2)
```

```
Requirement already satisfied: scipy>=1.3.2 in
```

```
/usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.11.4)
```

```
Requirement already satisfied: joblib>=1.1.1 in
```

```
/usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.4.2)
```

```
Requirement already satisfied: threadpoolctl>=2.0.0 in
```

```
/usr/local/lib/python3.10/dist-packages (from scikit-learn) (3.5.0)
```

```
classifier = KNeighborsClassifier(n_neighbors=5, metric='minkowski',
p=2)
```

```

classifier.fit(X_train, y_train)
KNeighborsClassifier()
y_pred = classifier.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
Accuracy: 0.93

from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_train, y_train)
LogisticRegression()
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
Accuracy: 0.89

from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X_train, y_train)
LinearRegression()
y_pred = model.predict(X_test)

from sklearn.metrics import mean_squared_error, r2_score
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print("Mean Squared Error:", mse)
print("R-squared:", r2)
Mean Squared Error: 0.09872045740502204
R-squared: 0.5463214273666266

from sklearn.tree import DecisionTreeClassifier
model = DecisionTreeClassifier()
model.fit(X_train, y_train)
DecisionTreeClassifier()
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
Accuracy: 0.91

```

```
from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier()
model.fit(X_train, y_train)

RandomForestClassifier()

y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

Accuracy: 0.93
```